

ESE 545 — Project Part 1

SPU-lite: Dual-Issue Processing Core with
Register File, Forwarding, and Write-Back

Name: Jing Jin

ID: 117502257

1 Introduction

This report presents the design and verification of **Part 1** of the SPU-lite processor, a dual-issue in-order pipeline modeled after the Synergistic Processing Unit (SPU) of the Cell Broadband Engine. Part 1 implements the back-end of the pipeline: the register file with read/write ports, operand forwarding, an 8-stage execution shift register, hazard detection with stalling, branch resolution with misprediction flush, and write-back.

The design supports all **102 instructions** defined in the SPU-lite ISA, organized across 8 execution units in two parallel pipes (even and odd). A comprehensive forwarding network eliminates most pipeline stalls by bypassing results directly from in-flight execution stages. A hardware interlock stalls the pipeline when a true data dependency cannot be resolved by forwarding. Branch instructions are resolved combinationally at the execution stage; on misprediction the pipeline is flushed and the correct target is redirected to the fetch unit.

Verification is carried out with four dedicated SystemVerilog testbenches that collectively demonstrate: (1) correct execution of all 102 instructions, (2) correct hazard detection and data forwarding across all unit latencies and cross-pipe scenarios, and (3) correct branch resolution for all 11 branch types including conditional taken/not-taken paths and link-register writes.

2 Module Architecture

2.1 Design Files Overview

The implementation consists of three synthesizable source files and a shared package, summarized in Table 1.

Table 1: Source files and their roles.

File	Lines	Description
<code>spu_pkg.sv</code>	228	Shared package: opcodes (0–102), <code>decoded_instr_t</code> and <code>result_pkt_t</code> structs, pipeline constants, word/half/byte slice helpers
<code>regfile.sv</code>	78	128×128-bit register file with 6 read ports, 2 write ports, and write-through bypass
<code>spu_core_part2.sv</code>	1072	Top-level execution core: forwarding, hazard detection, even/odd execution, branch resolution, pipe shift registers, write-back, local store

2.2 Pipeline Organization

The SPU-lite pipeline is an 11-stage in-order dual-issue design. Part 1 covers stages 3 through 11:

- **Stage 3 (RF/FWD):** The register file is read and the forwarding mux selects the most recent value for each source operand. Hazard detection runs in parallel; if a RAW hazard cannot be resolved by forwarding, the `stall_o` signal holds the upstream decode stage.
- **Stages 4–11 (EXEC/FWD/WB):** An 8-deep shift register (`even_pipe[0:7]` and `odd_pipe[0:7]`) carries `result_pkt_t` packets through execution. Each packet includes a `cycles_to_ready` counter that decrements every cycle; when it reaches zero the result is forwardable.

Two instructions are issued per cycle, one into each pipe:

- **Even pipe:** FX1 (depth 2), FX2 (depth 3), FP (depth 6), Integer Multiply (depth 7), Byte (depth 3).
- **Odd pipe:** Permute (depth 3), Load/Store (depth 6), Branch (depth 1).

Table 2 lists each unit with its latency and instruction count.

Table 2: Execution units.

Unit	Pipe	Depth	Latency	Count
Simple Fixed 1 (FX1)	Even	2	3	41
Simple Fixed 2 (FX2)	Even	3	4	10
SP Floating-Point (FP)	Even	6	7	11
Integer Multiply	Even	7	8	4
Byte	Even	3	4	10
Permute (PERM)	Odd	3	4	8
Load/Store (LS)	Odd	6	7	4
Branch (BR)	Odd	1	2	11
Control	—	—	—	3
Total				102

2.3 Register File (`regfile.sv`)

The register file contains 128 general-purpose 128-bit registers. It provides:

- **6 combinational read ports:** three per pipe (`ra/rb/rc` for each of the even and odd instructions).

- **2 synchronous write ports:** one for even-pipe write-back, one for odd-pipe write-back. If both ports target the same register in the same cycle, the even port takes priority.
- **Write-through bypass:** a read in the same cycle as a write to the same address returns the new value, avoiding a one-cycle stale-read hazard.

2.4 Forwarding Network

The `forward()` function scans all 8 stages of both the even and odd execution pipes (16 slots total), searching newest-first (index 0 = stage 4) for a matching destination register with `cycles_to_ready == 0`. If found, the in-flight result is returned; otherwise the register file value is used. This enables **cross-pipe forwarding**: an even-pipe result can be forwarded to an odd-pipe consumer and vice versa.

A special path handles **store-data forwarding**. For store instructions (STQD/STQA), the `rt` field identifies the data register (not a destination). The `b1` read port is repurposed to read `rt` when `is_store` is asserted, and the forwarding mux resolves `fw_opB_odd` from `rt` rather than `rb`. This ensures store-after-write hazards are correctly forwarded.

2.5 Hazard Detection

The `detect_stall()` function checks every source operand of both incoming instructions against all in-flight results. A stall is generated when:

1. An operand matches a result packet with `cycles_to_ready > 0` (the value is in the pipe but not yet forwardable).
2. An operand matches a Stage-3 instruction (`ex_even/ex_odd`) whose result has not yet entered the pipe.

On stall, a bubble (null instruction) is injected into Stage 3 while the upstream decode holds. Stall is suppressed during a flush cycle since the flushed instructions are discarded.

2.6 Branch Resolution

Branch instructions execute in the odd pipe. The `resolve_branch()` function computes the branch outcome combinationaly from `ex_odd`:

- **Unconditional:** BR/BRA are always taken. BRSL/BRASL are always taken and additionally write a link value (`PC + 4`) into `rt[127:96]`.
- **Conditional (word):** BRZ (taken if `word0(ra) == 0`), BRNZ (taken if `word0(ra) != 0`).
- **Conditional (halfword):** BRHZ (taken if `half0(ra) == 0`), BRHNZ (taken if `half0(ra) != 0`).
- **Indirect:** BI (always taken, target = `ra[10:0]`), BIZ/BINZ (conditional, target = `ra[10:0]`).

A **misprediction** is detected when the resolved outcome differs from the upstream prediction (`id_pred_taken_i`, `id_pred_target_i`). On misprediction, `flush_o` is asserted and `redirect_pc_o` emits the correct target (branch target if taken, `PC + 8` if not taken). The pipeline injects a bubble on flush, squashing the speculative pair.

3 Testbench Architecture

Four testbenches provide complementary coverage. Each builds instruction programs in memory arrays, drives them through the DUT, and compares actual register values or control signals against a software reference model. Table 3 summarizes the testbench suite.

Table 3: Testbench files and coverage.

Testbench	Coverage	Log
<code>spu_tb_instructions.sv</code>	All 102 opcodes with per-instruction PASS/FAIL	<code>instructions.log</code>
<code>spu_tb_hazard.sv</code>	61 hazard scenarios: RAW, WAW, forwarding chains, LS coherence	<code>hazard.log</code>
<code>spu_tb_branch.sv</code>	All 11 branch types, taken + not-taken, flush + link	<code>branch.log</code>
<code>spu_tb_rand.sv</code>	Random stress with mixed hazards (multi-seed)	<code>simulation_rand.log</code>

3.1 Instruction Coverage Testbench (`spu_tb_instructions.sv`)

This testbench systematically exercises every opcode. It uses a **hazard-free sweep** strategy: source registers `r1--r63` are preloaded with known values via `ILHU+ORI` pairs; every instruction under test writes a unique destination in `r64--r127`, ensuring no RAW dependencies. Three separate program runs keep destinations within the 64-register budget:

- **Run A:** 41 FX1 instructions (arithmetic, logic, compare, immediates).
- **Run B:** 34 FX2 + FP + IntMul + Byte instructions.
- **Run C:** 8 Permute + 4 Load/Store (store then load-back) + 2 branch-link (BRS�/BRASL) + 9 non-writing branches + NOP/LNOP/STOP.

A fourth run exercises **dual-issue**: even-pipe and odd-pipe instructions are placed in the *same pair* (e.g., FX1 + PERM, FX2 + PERM, BYTE + PERM, FP + STQD, FX1 + BRS�), proving both pipes execute real instructions simultaneously.

A full software reference model (`ref_exec`) mirrors every opcode. After each run, `check_all()` compares every destination register against the reference value and reports per-opcode PASS or FAIL.

3.2 Hazard and Forwarding Testbench (`spu_tb_hazard.sv`)

This testbench targets the forwarding, stall, and write-back logic with **61 individual test cases** spanning five hazard categories. Each test places instructions at exact pair indices with controlled gaps and verifies the final register value against a software reference model.

RAW (Read-After-Write) — 46 tests. Each test places a producer and a consumer at a controlled gap:

- **Same-pipe:** FX1 (gaps 0–3), FX2 (gaps 0–4), FP (gaps 0,1,5,6,7), IntMul (gaps 0,1,6,7,8), Byte (gaps 0–4), Permute (gaps 0–4), LS-load (gaps 0,1,5,6,7).
- **Store-data RAW:** Forwarding of store-data operand `rt` (gaps 0–3).
- **Cross-pipe:** Even→Odd (gaps 0–3), Odd→Even (gaps 0–3).

WAW (Write-After-Write) — 8 tests. Two instructions write the same destination register; the program-order-last write must win:

- **Same-unit WAW:** FX1→FX1 at gaps 0,1,2 — verifies the forwarding network returns the newest result.
- **Mixed-latency WAW:** FP (depth=6) then FX1 (depth=2) to the same register — FX1 is later in program order and must win despite FP being longer.
- **Cross-pipe WAW:** Even then Odd to the same register (gaps 0,1).

- **Same-pair WAW:** Even and Odd in the same pair write the same register — the register file even-port priority determines the winner.

Forwarding Chain — 2 tests. Multi-level RAW dependency: A produces a value consumed by B, whose result is consumed by C. Tests that the forwarding network correctly resolves cascaded dependencies:

- FX1→FX1→FX1 with tight gaps (stall + forward at each level).
- Cross-pipe chain: Even→Odd→Even.

Load-After-Store (LS coherence) — 4 tests. STQD writes to the local store, then LQD reads from the same address at gaps 1,2,3,7. Verifies that the store commits to the local store before the load reads it.

Same-pair even→odd RAW — 1 test. The even instruction writes a register that the odd instruction reads in the *same pair*. The core has no intra-pair forwarding; the odd instruction receives the old register value. This test documents the expected behavior.

3.3 Branch Testbench (spu_tb_branch.sv)

This testbench verifies all 11 branch instructions with **26 test cases**: 18 basic tests (taken/not-taken for each conditional type) plus 8 branch-prediction tests that exercise all four prediction outcomes. The strategy:

- A fork-based monitor captures `flush_o` and `redirect_pc_o` concurrently with program driving.
- For unconditional branches (BR, BRA): verify `flush_o` and correct target.
- For BRSL/BRASL: additionally verify the link register `rt[127:96] = {21'd0, PC+4}` with the remaining 96 bits all zero.
- For conditional branches: use `r0=0` (from reset) as the “zero” condition register and a preloaded `r1=0x00010000` as the “nonzero” register; test both taken and not-taken paths.
- For indirect branches (BI, BIZ, BINZ): preload `r5=0x70` as the target address; verify the redirect points to `0x070`.
- **Prediction tests:** The DUT accepts a branch prediction from upstream (`id_pred_taken_i`, `id_pred_target_i`) and compares it against the resolved outcome. The testbench exercises all four prediction×outcome combinations:
 1. `pred=0`, not-taken → correct prediction, no flush
 2. `pred=0`, taken → direction mismatch, flush
 3. `pred=1` + correct target, taken → correct prediction, no flush
 4. `pred=1` + wrong target, taken → target mismatch, flush
 5. `pred=1`, not-taken → direction mismatch, flush to PC+8

3.4 Random Stress Testbench (spu_tb_rand.sv)

This testbench runs multi-phase random programs with multiple random seeds:

- **Phase 1:** Hazard-free sweep of all 102 opcodes to establish a baseline.
- **Phase 2:** Deliberate hazard stress — tightly scheduled producer/consumer pairs with minimal gaps to exercise forwarding and stalling.
- **Phase 3:** Fully random instruction sequences with random register assignments, generating both hazards and hazard-free pairs naturally.

After each phase, every written register is compared against the software reference model. The testbench runs with two different seeds to increase coverage. Timing annotations (`$time`) mark each phase boundary in the log.

4 Functionality Demonstration

4.1 All 102 Instructions Successfully Implemented

The instruction coverage testbench (`spu_tb_instructions.sv`) verifies every opcode in the ISA by emitting one instance of each instruction and comparing the destination register against a bit-exact software reference model. The three runs cover the full instruction set:

- **Run A (41 FX1 instructions):** IL, ILH, ILHU, ILA, A, AI, SF, SFI, AH, AHI, SFH, SFHI, CG, BG, ADDX, SFX, AND, OR, XOR, NOR, NAND, EQV, ANDI, ORI, XORI, ANDHI, ORHI, XORHI, CEQ, CEQI, CGT, CGTI, DGT, DGTI, CEQH, CEQHI, CGTHI, DGTH, DGTHI, CLZ, SELB.
- **Run B (34 FX2+FP+IntMul+Byte instructions):** SHL, SHLI, SHLH, SHLHI, SHRA, SHRAI, ROT, ROTI, ROTH, ROTH, FA, FS, FM, FMA, FMS, FMS2, FNMS, FCEQ, FCGT, FCMEQ, FCMGT, MPY, MPYU, MPYA, MPYI, ABSDB, AVGB, SUMB, CNTB, FSMB, FSMBI, CEQB, CGTB, CLGTB, GB.
- **Run C (27 PERM+LS+BR+Control instructions):** SHLQBI, SHLQBII, SHLQBY, SHLQBYI, ROTQBY, ROTQBYI, ROTQBI, ROTQBII, STQD+LQD, STQA+LQA, BRSL, BRASL, BR, BRA, BRZ, BRNZ, BRHZ, BRHNZ, BI, BIZ, BINZ, NOP, LNOP, STOP.
- **Run D (Dual-issue, 11 checks):** Even-pipe and odd-pipe instructions execute in the *same pair*: A+ROTQBYI, SHLI+SHLQBYI, AVGB+ROTQBYI, FA+STQD (with LQD verify), CEQI+BRSL (with link check). Proves both pipes are active simultaneously.

Since source registers are preloaded with deterministic values and destinations are unique per instruction, any failure is guaranteed to be a pure compute error rather than a pipeline hazard artifact. Expected output in `instructions.log`:

```
PASS: 102    FAIL: 0    ALL 102 INSTRUCTIONS PASSED
```

```

— Run A — FX1 —
PASS IL          r64 = 0xffffffffffffffffffffffffffffffff
PASS ILH         r65 = 0x12341234123412341234123412341234
PASS ILHU        r66 = 0xdead0000dead0000dead0000dead0000
PASS ILA         r67 = 0x0003cafe0003cafe0003cafe0003cafe
PASS A           r68 = 0xd2610571d2610571d2610571d2610571
PASS AI          r69 = 0x640b02c6640b02c6640b02c6640b02c6
PASS SF          r70 = 0x0a4b002f0a4b002f0a4b002f0a4b002f
PASS SFI         r71 = 0x9bf4fdc29bf4fdc29bf4fdc29bf4fdc2
PASS AH          r72 = 0xd2610571d2610571d2610571d2610571
PASS AHI         r73 = 0x641202a8641202a8641202a8641202a8
PASS SFH         r74 = 0x0a4b002f0a4b002f0a4b002f0a4b002f
PASS SFHI        r75 = 0x9bf8fd629bf8fd629bf8fd629bf8fd62
PASS CG          r76 = 0x00000000000000000000000000000000
PASS BG          r77 = 0x00000001000000010000000100000001
PASS ADDX        r78 = 0xd2610571d2610571d2610571d2610571
PASS SFX         r79 = 0x0a4b002e0a4b002e0a4b002e0a4b002e
PASS AND         r80 = 0x64020280640202806402028064020280
PASS OR          r81 = 0x6e5f02f16e5f02f16e5f02f16e5f02f1
PASS XOR         r82 = 0x0a5d00710a5d00710a5d00710a5d0071
PASS NOR         r83 = 0x91a0fd0e91a0fd0e91a0fd0e91a0fd0e
PASS NAND        r84 = 0x9bdfd7f9bdfd7f9bdfd7f9bdfd7f9bdfd7f
PASS EQV         r85 = 0xf5a2ff8ef5a2ff8ef5a2ff8ef5a2ff8e
PASS ANDI        r86 = 0x00000000000000000000000000000000
PASS ORI         r87 = 0x640b02a5640b02a5640b02a5640b02a5
PASS XORI        r88 = 0x640b025e640b025e640b025e640b025e
PASS ANDHI       r89 = 0x000b0001000b0001000b0001000b0001
PASS ORHI        r90 = 0x64fb02f164fb02f164fb02f164fb02f1
PASS XORHI       r91 = 0x64a1020b64a1020b64a1020b64a1020b
PASS CEQ         r92 = 0x00000000000000000000000000000000
PASS CEQI        r93 = 0x00000000000000000000000000000000
PASS CGT         r94 = 0x00000000000000000000000000000000
PASS CGTI        r95 = 0xffffffffffffffffffffffffffffffff
PASS DGT         r96 = 0x00000000000000000000000000000000
PASS DGTI        r97 = 0xffffffffffffffffffffffffffffffff
PASS CEQH        r98 = 0x00000000000000000000000000000000
PASS CEQHI       r99 = 0x00000000000000000000000000000000
PASS CGTHI       r100 = 0xffffffffffffffffffffffffffffffff
PASS DGTH        r101 = 0x00000000000000000000000000000000
PASS DGTHI       r102 = 0xffffffffffffffffffffffffffffffff
PASS CLZ         r103 = 0x00000001000000010000000100000001
PASS SELB        r104 = 0x6c0702d16c0702d16c0702d16c0702d1
[t=3185000 ns] Run B start (244 pairs, 35 tests)

```

Figure 1: Simulation log showing 41 FX1 instructions

```

[t=3185000 ns] Run B start (244 pairs, 35 tests)
— Run B — FX2+FP+IntMul+BYTE —
PASS SHL      r64 = 0xd6000000d6000000d6000000d6000000
PASS SHLI     r65 = 0xf2803580f2803580f2803580f2803580
PASS SHLH     r66 = 0x2800d6002800d6002800d6002800d600
PASS SHLHI    r67 = 0x3ca00d603ca00d603ca00d603ca00d60
PASS SHRA     r68 = 0x00000008f0000008f0000008f0000008f
PASS SHRAI    r69 = 0x0479401a0479401a0479401a0479401a
PASS ROT      r70 = 0xd623ca00d623ca00d623ca00d623ca00
PASS ROTI     r71 = 0xca00d623ca00d623ca00d623ca00d623
PASS ROTH     r72 = 0x288fd600288fd600288fd600288fd600
PASS ROTHI    r73 = 0xf2883580f2883580f2883580f2883580
PASS FA       r74 = 0xfffd90245fffd90245fffd90245fffd90245
PASS FS       r75 = 0xfffd90245fffd90245fffd90245fffd90245
PASS FM       r76 = 0xfffd90245fffd90245fffd90245fffd90245
PASS FMA      r77 = 0xfffd90245fffd90245fffd90245fffd90245
PASS FMS      r78 = 0xfffd90245fffd90245fffd90245fffd90245
PASS FMS2     r79 = 0xfffd90245fffd90245fffd90245fffd90245
PASS FNMS     r80 = 0x7fd902457fd902457fd902457fd90245
PASS FCEQ     r81 = 0x00000000000000000000000000000000
PASS FCGT     r82 = 0x00000000000000000000000000000000
PASS FCMEQ    r83 = 0x00000000000000000000000000000000
PASS FCMGT    r84 = 0x00000000000000000000000000000000
PASS MPY      r85 = 0x00025ee400025ee400025ee400025ee4
PASS MPYU     r86 = 0x00025ee400025ee400025ee400025ee4
PASS MPYA     r87 = 0x682c6171682c6171682c6171682c6171
PASS MPYI     r88 = 0x000015bc000015bc000015bc000015bc
PASS ABSDB    r89 = 0x70630041706300417063004170630041
PASS AVGB     r90 = 0x7f63018c7f63018c7f63018c7f63018c
PASS SUMB     r91 = 0x01c3011901c3011901c3011901c30119
PASS CNTB     r92 = 0x04030104040301040403010404030104
PASS FSMB     r93 = 0xffffffffffffffffffffffffffffffff
PASS FSMBI    r94 = 0xfffffffff0000000fffffffff0000000
PASS CEQB     r95 = 0x0000ff000000ff000000ff000000ff00
PASS CGTB     r96 = 0xff000000ff000000ff000000ff000000
PASS CLGTB    r97 = 0x00ff00ff00ff00ff00ff00ff00ff00ff
PASS GB       r98 = 0x00005555000055550000555500005555

```

Figure 2: Simulation log showing 34 FX2+FP+IntMul+Byte instructions

```

[t=6315000 ns] Run C start (232 pairs, 11 tests)
— Run C — PERM+LS+BRSL —
PASS SHLQBI   r64 = 0x77801edc77801edc77801edc77801ed8
PASS SHLQBII  r65 = 0x3bc00f6e3bc00f6e3bc00f6e3bc00f6c
PASS SHLQBY   r66 = 0x00000000000000000000000000000000
PASS SHLQBYI  r67 = 0xdb8ef003db8ef003db8ef003db000000
PASS ROTQBY   r68 = 0xdb8ef003db8ef003db8ef003db8ef003
PASS ROTQBYI  r69 = 0xf003db8ef003db8ef003db8ef003db8e
PASS ROTQBII  r70 = 0x77801edc77801edc77801edc77801edc
PASS ROTQBII  r71 = 0x77801edc77801edc77801edc77801edc
PASS STQD+LQD r72 = 0x8ef003db8ef003db8ef003db8ef003db
PASS STQA+LQA r73 = 0x836e00eb836e00eb836e00eb836e00eb
PASS BRSL     r74 = 0x0000069c000000000000000000000000
[t=9325000 ns] Run D start (210 pairs, 1 tests)
— Run D — BRASL —
PASS BRASL    r64 = 0x0000063c000000000000000000000000

```

Figure 3: Simulation log showing 27 PERM+LS+BR+Control instructions

```

7 — Run D — Dual-Issue (even+odd same cycle) —————
5 PASS Dual A(even) r64 = 0x5c7504e95c7504e95c7504e95c7504e9
4 PASS Dual ROTQBYI(odd) r65 = 0x00885b6e00885b6e00885b6e00885b6e
3 PASS Dual SHLI(even) r66 = 0x2bf028002bf028002bf028002bf02800
2 PASS Dual SHLQBYI(odd) r67 = 0x6959b6026959b6026959b60269000000
1 PASS Dual AVGB(even) r68 = 0x2ebb02752ebb02752ebb02752ebb0275
0 PASS Dual ROTQBYI(odd) r69 = 0x4103e9b84103e9b84103e9b84103e9b8
9 PASS Dual FA(even)+STQD(odd) r70 = 0xcdcf0312cdf0312cdf0312cdf0312
8 PASS Dual verify LQD(odd) r71 = 0x02bf028002bf028002bf028002bf0280
7 PASS Dual CEQI(even) r72 = 0x00000000000000000000000000000000
6 PASS Dual BRSL-Ltnk(odd) r73 = 0x000006a4000000000000000000000000

```

Figure 4: Simulation log showing dual-issue




Figure 5: Simulation waveform showing dual-issue

4.2 Hazard Detection and Data Forwarding

The hazard testbench (`spu_tb_hazard.sv`) stresses the forwarding network, stall logic, and write-back arbitration with 61 targeted scenarios across five hazard categories:

1. **RAW — Forwarding at exact depth boundary (46 tests):** For each execution unit, a producer/consumer pair is placed at gaps spanning from 0 (maximum stall) through the forwarding boundary to a gap beyond the depth (no stall). Cross-pipe forwarding (even→odd, odd→even) and store-data forwarding via the repurposed `b1` port are also covered.
2. **WAW — Write ordering (8 tests):** Two instructions writing the same register verify that the program-order-last write wins. Cases include same-unit (FX1→FX1), mixed-latency (FP→FX1 where the shorter-latency FX1 is later), cross-pipe (even→odd), and same-pair (even-port priority).
3. **Forwarding chain (2 tests):** A→B→C multi-level RAW dependencies with tight gaps validate that cascaded forwarding resolves correctly, both within a single pipe and across pipes.
4. **Load-after-store coherence (4 tests):** A STQD followed by a LQD to the same local store address at gaps 1,2,3,7 verifies that the store commits before the load reads.
5. **Same-pair intra-pair RAW (1 test):** Even writes a register that odd reads in the same pair. Since no intra-pair forwarding exists, odd receives the old value — the test verifies this documented behavior.

Expected output in `hazard.log`:

```
PASS: 61   FAIL: 0   ALL HAZARD TESTS PASSED
```


Table 4: Branch direction/target test cases.

Instruction	Type	Tests	Taken Condition	Target
BR	Unconditional	1	Always	PC-relative
BRA	Unconditional	1	Always	Absolute
BRSL	Uncond. + link	1	Always	PC-relative; link = {21'd0, PC+4}
BRASL	Uncond. + link	1	Always	Absolute; link = {21'd0, PC+4}
BRZ	Cond. word	2	word0(ra)==0	PC-relative
BRNZ	Cond. word	2	word0(ra)!=0	PC-relative
BRHZ	Cond. half	2	half0(ra)==0	PC-relative
BRHNZ	Cond. half	2	half0(ra)!=0	PC-relative
BI	Indirect	1	Always	ra[10:0]
BIZ	Indirect cond.	2	word0(ra)==0	ra[10:0]
BINZ	Indirect cond.	2	word0(ra)!=0	ra[10:0]
Subtotal		18		

For BRSL and BRASL, the testbench additionally reads the destination register from the register file via hierarchical access (`dut.u.rf.regs[rt]`) and verifies:

- `word0` = {21'd0, `branch_PC` + 4} (return address in word 0).
- `word1` = `word2` = `word3` = 0 (remaining words are zero).

Group 2: Branch prediction verification (8 tests). These tests exercise the misprediction detection logic by varying `id_pred_taken_i` and `id_pred_target_i`. The DUT detects a misprediction when: `(taken ≠ pred_taken) || (taken && target ≠ pred_target)`. On misprediction, `flush_o=1` and `redirect_pc_o` emits the correct target (branch target if taken, PC+8 if not taken). Table 5 shows the prediction test matrix.

Table 5: Branch prediction test cases.

Branch	Prediction	Actual	Scenario	Flush?
BR	pred=1, correct target	taken	Correct prediction	No
BR	pred=1, wrong target	taken	Target mismatch	Yes
BRZ	pred=1	not-taken (r1≠0)	Direction mismatch	Yes
BRNZ	pred=0	not-taken (r0=0)	Correct prediction	No
BRA	pred=1, correct target	taken	Correct prediction	No
BRSL	pred=1, correct target	taken	Correct pred. + link check	No
BI	pred=1, correct target	taken	Correct prediction	No
BI	pred=1, wrong target	taken	Target mismatch	Yes
Subtotal				8

Expected output in `branch.log`:

PASS: 26 FAIL: 0 ALL BRANCH TESTS PASSED

```

24 SPU-lite - Branch Testbench t=0 ns
23 P BR=6 BR_PC=0x030 BR_REL_TGT=0x050 BRA_TGT=0x0f0
22 BRSL_LINK_W0=0x00000034 BI_TGT=0x070
21
20
19 — Unconditional branches
18 PASS BR (taken) flush=1 target=0x050
17 PASS BRA (taken) flush=1 target=0x0f0
16 PASS BRSL (taken + link) flush=1 target=0x050 link_w0=0x00000034
15 PASS BRASL (taken + link) flush=1 target=0x0f0 link_w0=0x00000034
14 — BRZ / BRNZ
13 PASS BRZ taken (ra=r0=0) flush=1 target=0x050
12 PASS BRZ not-taken (ra=r1≠0) flush=0 target=0x000
11 PASS BRNZ taken (ra=r1≠0) flush=1 target=0x050
10 PASS BRNZ not-taken (ra=r0=0) flush=0 target=0x000
9 — BRHZ / BRHNZ
8 PASS BRHZ taken (ra=r0, half0=0) flush=1 target=0x050
7 PASS BRHZ not-taken (ra=r1, half0=0x0001) flush=0 target=0x000
6 PASS BRHNZ taken (ra=r1, half0=0x0001) flush=1 target=0x050
5 PASS BRHNZ not-taken (ra=r0, half0=0) flush=0 target=0x000
4 — BI / BIZ / BINZ
3 PASS BI taken (r5=0x070) flush=1 target=0x070
2 PASS BIZ taken (ra=r0=0, tgt=0) flush=1 target=0x000
1 PASS BIZ not-taken (ra=r5≠0) flush=0 target=0x000
15 PASS BINZ taken (ra=r5≠0, tgt=0x070) flush=1 target=0x070
1 PASS BINZ not-taken (ra=r0=0) flush=0 target=0x000
2 — Branch prediction
3 PASS BR pred=1 correct tgt → no flush flush=0 target=0x000
4 PASS BR pred=1 wrong tgt → flush flush=1 target=0x050
5 PASS BRZ pred=1 not-taken → flush flush=1 target=0x038
6 PASS BRNZ pred=0 not-taken → no flush flush=0 target=0x000
7 PASS BRA pred=1 correct tgt → no flush flush=0 target=0x000
8 PASS BRSL pred=1 correct → no flush+link flush=0 target=0x000 link_w0=0x00000034
9 PASS BI pred=1 correct tgt → no flush flush=0 target=0x000
10 PASS BI pred=1 wrong tgt → flush flush=1 target=0x070
11
12 PASS: 25 FAIL: 0 [t=6000000 ns]
13 ALL BRANCH TESTS PASSED
14

```

Figure 8: Simulation log showing branch




Figure 9: Simulation waveform showing flush is set when branch mispredicted

5 Conclusion

The SPU-lite Part 1 core implements the full register file, operand forwarding, 8-stage execution pipeline, hazard detection, branch resolution, and write-back for a dual-issue 102-instruction ISA. Four dedicated testbenches provide comprehensive verification:

- **Instruction coverage:** all 102 opcodes produce bit-exact results.
- **Hazard/forwarding:** 61 scenarios across five hazard types (RAW, WAW, forwarding chains, LS coherence, intra-pair) confirm correct stalling, forwarding, and write-back arbitration.
- **Branch resolution:** all 11 branch types correctly compute flush and redirect signals for both taken and not-taken outcomes; BRSL/BASL link registers are verified.
- **Random stress:** multi-seed random programs with mixed hazards confirm end-to-end correctness under realistic workloads.